

week12-01-least-squares

April 22, 2024

1 George McNinch Math 87 - Spring 2024

2 Week 12

3 Least squares

4 Linear Least Squares

We are going to begin our discussion of “least squares” approximation with an example.

4.1 Example

Consider a stretch of highway with four distinct reference points **A**, **B**, **C**, and **D**:

[**A**] — x_1 — [**B**] — x_2 — [**C**] — x_3 — [**D**]

Write $x_1 = \mathbf{AB}$ for the distance from **A** to **B**, $x_2 = \mathbf{BC}$, $x_3 = \mathbf{CD}$.

We take some measurements – which potentially reflect errors – , and we seek the *best approximation* to the distances x_1 , x_2 , x_3 .

The measurements taken are as follows:

segment	AD	AC	BD	AB	CD
length	89 m	67 m	53 m	35 m	20 m

Thus the observations suggest the following equations:

$$(1) \quad x_1 + x_2 + x_3 = 89$$

$$(2) \quad x_1 + x_2 = 67$$

$$(3) \quad x_2 + x_3 = 53$$

$$(4) \quad x_1 = 35$$

$$(5) \quad x_3 = 20$$

These equations aren’t compatible, though. Note e.g. that equations (3) -- (5) indicate the following:

$$x_1 = 35, \quad x_3 = 20, \quad x_2 = 53 - 20 = 33$$

but then we find that

$$x_1 + x_2 + x_3 = 35 + 33 + 20 = 88$$

which is incompatible with (1).

And we find that

$$x_1 + x_2 = 35 + 33 = 68$$

which is incompatible with (2).

Let's formulate these equalities in matrix form.

Thus let

$$A = \begin{pmatrix} 1 & 1 & 1 \\ 1 & 1 & 0 \\ 0 & 1 & 1 \\ 1 & 0 & 0 \\ 0 & 0 & 1 \end{pmatrix}, \quad \mathbf{x} = \begin{pmatrix} x_1 \\ x_2 \\ x_3 \end{pmatrix}, \quad \mathbf{b} = \begin{pmatrix} 89 \\ 67 \\ 53 \\ 35 \\ 20 \end{pmatrix}.$$

With these notations, the above equations suggest that $A\mathbf{x}$ should be equal to \mathbf{b} .

Our observation(s) in the preceding slides show, however, that the system of equations $A\mathbf{x} = \mathbf{b}$ is *inconsistent* (i.e. there is no vector \mathbf{x} which makes the equation true).

4.2 Residual

In general, given an $m \times n$ matrix A , a column vector $\mathbf{b} \in \mathbb{R}^m$ and an equation $A\mathbf{x} = \mathbf{b}$, we instead look at the so-called *residual*

$$\mathbf{r} = \mathbf{b} - A\mathbf{x},$$

and *minimize* this residual.

More precisely, we want to minimize the *magnitude* (or *length*) of this vector.

Thus if $\mathbf{r} = (r_1 \ \cdots \ r_m)^T$, we must minimize the quantity

$$\|\mathbf{r}\| = \left(\sum_{i=1}^m r_i^2 \right)^{1/2}$$

Here $\|\mathbf{r}\|$ is the magnitude, also called the Euclidean norm, of the vector \mathbf{r} .

In fact, because $f(x) = \sqrt{x}$ is an increasing function of x , we instead minimize the *square* of the magnitude of \mathbf{r} :

$$\|\mathbf{r}\|^2 = \sum_{i=1}^m r_i^2$$

Thus, we wish to find

$$\min_{\mathbf{x}} \|\mathbf{r}\|^2 = \min_{\mathbf{x}} \|\mathbf{b} - A\mathbf{x}\|^2 = \min_{x_1, x_2, \dots, x_n} \sum_{i=1}^m \left(b_i - \sum_{j=1}^n A_{ij} x_j \right)^2$$

The idea behind this minimization is to first compute for $1 \leq k \leq n$ the partial derivatives $\frac{\partial F}{\partial x_k}$ of the function

$$F(x_1, x_2, \dots, x_n) = \sum_{i=1}^m \left(b_i - \sum_{j=1}^n A_{ij} x_j \right)^2$$

Critical points - and thus possible minima - for F occur at points \mathbf{x} for which all $\frac{\partial F}{\partial x_k}(\mathbf{x}) = 0$.

Now,

$$\frac{\partial F}{\partial x_k} = \sum_{i=1}^m 2 \left(b_i - \sum_{j=1}^n A_{ij} x_j \right) (-A_{ik}) = 2 \left(\sum_{i=1}^m (-A_{ik} b_i) + \sum_{i=1}^m A_{ik} \sum_{j=1}^n A_{ij} x_j \right)$$

and this expression is equal to the k -th coefficient of the vector

$$2(-A^T \mathbf{b} + A^T A \mathbf{x})$$

Thus, the condition $\frac{\partial F}{\partial x_k} = 0$ for all k is equivalent to the so-called *normal equations*:

$$(\diamond) \quad A^T A \mathbf{x} = A^T \mathbf{b}.$$

Thus the solutions \mathbf{x} to the normal equations (\diamond) are precisely the critical points of the function F .

Recall that $A \in \mathbb{R}^{m \times n}$. Thus, $A^T \in \mathbb{R}^{n \times m}$ so that the matrix $A^T A$ is $n \times n$; in particular, $A^T A$ is always a square matrix.

Moreover, $A^T A$ is *symmetric*, since

$$(A^T A)^T = A^T (A^T)^T = A^T A.$$

We are interested here in the case of *overdetermined systems* – i.e. in the case where A has more rows than columns (“more equations than variables”). Thus $m \geq n$.

We also are interested in the case where A has rank n – i.e. A has n linearly independent columns – since otherwise we don’t expect to have enough information to find \mathbf{x} .

4.3 Proposition

Let $A \in \mathbb{R}^{m \times n}$, suppose that $m \geq n$ and that A has rank n . Then $A^T \cdot A$ is invertible.

Proof:

Since $A^T \cdot A$ is an $n \times n$ square matrix, the proposition will follow if we argue that the null space $\text{Null}(A^T A)$ is zero. So: suppose that $\mathbf{v} \in \text{Null}(A^T A)$.

Thus $A^T A \mathbf{v} = 0$ and thus also $\mathbf{v}^T A^T \cdot A \mathbf{v} = 0$.

Now,

$$0 = \mathbf{v}^T A^T \cdot A \mathbf{v} = (A \mathbf{v})^T (A \mathbf{v})$$

and of course for any vector \mathbf{w} , we know that

$$0 = \mathbf{w}^T \mathbf{w} \implies \mathbf{w} = \mathbf{0}.$$

We now conclude that $A\mathbf{v} = 0$, so $\mathbf{v} \in \text{Null}(A)$. Since A has rank n , the Null space of A is equal to zero, and we conclude that $\mathbf{v} = \mathbf{0}$.

We have now proved that $\text{Null}(A^T A)$ is zero, as required.

Remark: What we have really showed is that the symmetric matrix $A^T A$ is *definite*: $\mathbf{v}^T A^T A \mathbf{v} = 0 \implies \mathbf{v} = \mathbf{0}$.

We finally claim that this solution must minimize the magnitude of the residual.

This depends on a “second derivative test” argument for which I’m not going to give full details. The main point is that the “second derivative” in this context – known as the Hessian – coincides with the matrix $2A^T A$. Now, under our assumptions the matrix $A^T A$ is positive definite, and it follows that \mathbf{x}_0 is a global minimum for the magnitude of the residual!

4.4 Return to the example

Recall that

$$A = \begin{pmatrix} 1 & 1 & 1 \\ 1 & 1 & 0 \\ 0 & 1 & 1 \\ 1 & 0 & 0 \\ 0 & 0 & 1 \end{pmatrix}, \quad \mathbf{x} = \begin{pmatrix} x_1 \\ x_2 \\ x_3 \end{pmatrix}, \quad \mathbf{b} = \begin{pmatrix} 89 \\ 67 \\ 53 \\ 35 \\ 20 \end{pmatrix}.$$

So to minimize the magnitude of the residual, we must solve the normal equations:

$$A^T A \mathbf{x} = A^T \mathbf{b}.$$

Now

$$A^T A = \begin{pmatrix} 1 & 1 & 0 & 1 & 0 \\ 1 & 1 & 1 & 0 & 0 \\ 1 & 0 & 1 & 0 & 1 \end{pmatrix} \cdot \begin{pmatrix} 1 & 1 & 1 \\ 1 & 1 & 0 \\ 0 & 1 & 1 \\ 1 & 0 & 0 \\ 0 & 0 & 1 \end{pmatrix} = \begin{pmatrix} 3 & 2 & 1 \\ 2 & 3 & 2 \\ 1 & 2 & 3 \end{pmatrix}$$

and

$$A^T \mathbf{b} = \begin{pmatrix} 191 \\ 209 \\ 162 \end{pmatrix}$$

So we need to solve the equation

$$\begin{pmatrix} 3 & 2 & 1 \\ 2 & 3 & 2 \\ 1 & 2 & 3 \end{pmatrix} \cdot \begin{pmatrix} x_1 \\ x_2 \\ x_3 \end{pmatrix} = \begin{pmatrix} 191 \\ 209 \\ 162 \end{pmatrix}$$

```
[1]: import numpy as np
import numpy.linalg as la

A= np.array([[1,1,1],[1,1,0],[0,1,1],[1,0,0],[0,0,1]])
b = np.array([89,67,53,35,20])

x0=la.solve(A.T @ A, A.T @ b)

def residual(x):
    return b - A @ x

def magnitude(x):
    return np.sqrt(x@x)

[x0,magnitude(residual(x0))]
```

```
[1]: [array([35.125, 32.5  , 20.625]), 1.1726039399558574]
```

Thus the *least squares solution* is

$$\mathbf{x}_0 = \begin{pmatrix} 35.125 \\ 32.5 \\ 20.625 \end{pmatrix}$$

and

$$\|\mathbf{b} - A\mathbf{x}_0\| \approx 1.1726$$

Recall that our “first guess” for a solution (based on some of the measurements) was

$$\mathbf{x}_1 = \begin{pmatrix} 35 \\ 33 \\ 20 \end{pmatrix}$$

The residual is indeed larger for x_1 :

```
[2]: x1 = np.array([35,33,20])

magnitude(residual(x1))
```

```
[2]: 1.4142135623730951
```

Let’s note that `numpy` already implements this least-squares functionality:

– you can [read more about it here](#)

```
[3]: res=la.lstsq(A,b,rcond=None)
res[0]
```

```
[3]: array([35.125, 32.5 , 20.625])
```

5 Example recapitulated

Consider towns in a certain area labeled [a,b,c,d,e,f,g,h,i,j]. You have hired someone to estimate the population of these towns, but they got confused and have reported the populations of *pairs* of towns.

Thus you have the following data:

```
[4]: p_est = {('a', 'b'): 21.21,
('a', 'c'): 24.35,
('a', 'd'): 34.57,
('a', 'e'): 32.72,
('a', 'f'): 36.14,
('a', 'g'): 17.21,
('a', 'h'): 26.41,
('a', 'i'): 29.92,
('a', 'j'): 32.84,
('b', 'c'): 29.09,
('b', 'd'): 33.78,
('b', 'e'): 40.63,
('b', 'f'): 44.46,
('b', 'g'): 17.03,
('b', 'h'): 30.58,
('b', 'i'): 32.7,
('b', 'j'): 31.19,
('c', 'd'): 30.08,
('c', 'e'): 37.3,
('c', 'f'): 39.4,
('c', 'g'): 20.44,
('c', 'h'): 28.34,
('c', 'i'): 31.07,
('c', 'j'): 36.78,
('d', 'e'): 41.44,
('d', 'f'): 47.97,
('d', 'g'): 28.51,
('d', 'h'): 38.0,
('d', 'i'): 38.71,
('d', 'j'): 39.24,
('e', 'f'): 59.61,
('e', 'g'): 29.19,
('e', 'h'): 35.09,
('e', 'i'): 42.18,
```

```
('e', 'j'): 46.8,  
('f', 'g'): 34.14,  
('f', 'h'): 46.71,  
('f', 'i'): 48.07,  
('f', 'j'): 49.53,  
('g', 'h'): 23.46,  
('g', 'i'): 22.13,  
('g', 'j'): 24.21,  
('h', 'i'): 29.03,  
('h', 'j'): 32.51,  
('i', 'j'): 38.17}
```

p_est

```
[4]: {('a', 'b'): 21.21,  
      ('a', 'c'): 24.35,  
      ('a', 'd'): 34.57,  
      ('a', 'e'): 32.72,  
      ('a', 'f'): 36.14,  
      ('a', 'g'): 17.21,  
      ('a', 'h'): 26.41,  
      ('a', 'i'): 29.92,  
      ('a', 'j'): 32.84,  
      ('b', 'c'): 29.09,  
      ('b', 'd'): 33.78,  
      ('b', 'e'): 40.63,  
      ('b', 'f'): 44.46,  
      ('b', 'g'): 17.03,  
      ('b', 'h'): 30.58,  
      ('b', 'i'): 32.7,  
      ('b', 'j'): 31.19,  
      ('c', 'd'): 30.08,  
      ('c', 'e'): 37.3,  
      ('c', 'f'): 39.4,  
      ('c', 'g'): 20.44,  
      ('c', 'h'): 28.34,  
      ('c', 'i'): 31.07,  
      ('c', 'j'): 36.78,  
      ('d', 'e'): 41.44,  
      ('d', 'f'): 47.97,  
      ('d', 'g'): 28.51,  
      ('d', 'h'): 38.0,  
      ('d', 'i'): 38.71,  
      ('d', 'j'): 39.24,  
      ('e', 'f'): 59.61,  
      ('e', 'g'): 29.19,  
      ('e', 'h'): 35.09,
```

```

('e', 'i'): 42.18,
('e', 'j'): 46.8,
('f', 'g'): 34.14,
('f', 'h'): 46.71,
('f', 'i'): 48.07,
('f', 'j'): 49.53,
('g', 'h'): 23.46,
('g', 'i'): 22.13,
('g', 'j'): 24.21,
('h', 'i'): 29.03,
('h', 'j'): 32.51,
('i', 'j'): 38.17}

```

The value `mdist[('a','b')] == 21.21` means that the sum of the populations of towns a and b is 21.21 thousand people.

Let's form the matrix `M` and vector `b` so that `M @ x == b` expresses these distance relationships.

```

[7]: towns = [ 'a', 'b', 'c', 'd', 'e', 'f', 'g', 'h', 'i', 'j' ]
      pairs = list(p_est.keys())

```

```

[8]: def sbv(i,n):
      # return the ith standard basis vector of length n
      return np.array([1 if j == i else 0 for j in range(n)])

      def sbv_list(elem,ls):
          # return the standard basis vector determined by the position of `elem` in
          ↪ the list `ls`
          return sbv(list(ls).index(elem),len(ls))

      M = np.array([sbv_list(x,towns) + sbv_list(y,towns) for (x,y) in p_est.keys()])

      b = np.array([p_est[x] for x in p_est.keys()])

      (M,b)

```

```

[8]: (array([[1, 1, 0, 0, 0, 0, 0, 0, 0, 0],
             [1, 0, 1, 0, 0, 0, 0, 0, 0, 0],
             [1, 0, 0, 1, 0, 0, 0, 0, 0, 0],
             [1, 0, 0, 0, 1, 0, 0, 0, 0, 0],
             [1, 0, 0, 0, 0, 1, 0, 0, 0, 0],
             [1, 0, 0, 0, 0, 0, 1, 0, 0, 0],
             [1, 0, 0, 0, 0, 0, 0, 1, 0, 0],
             [1, 0, 0, 0, 0, 0, 0, 0, 1, 0],
             [1, 0, 0, 0, 0, 0, 0, 0, 0, 1],
             [0, 1, 1, 0, 0, 0, 0, 0, 0, 0],
             [0, 1, 0, 1, 0, 0, 0, 0, 0, 0],
             [0, 1, 0, 0, 1, 0, 0, 0, 0, 0],

```



```

[0, 1, 0, 0, 0, 1, 0, 0, 0, 0],
[0, 1, 0, 0, 0, 0, 1, 0, 0, 0],
[0, 1, 0, 0, 0, 0, 0, 1, 0, 0],
[0, 1, 0, 0, 0, 0, 0, 0, 1, 0],
[0, 1, 0, 0, 0, 0, 0, 0, 0, 1],
[0, 0, 1, 1, 0, 0, 0, 0, 0, 0],
[0, 0, 1, 0, 1, 0, 0, 0, 0, 0],
[0, 0, 1, 0, 0, 1, 0, 0, 0, 0],
[0, 0, 1, 0, 0, 0, 1, 0, 0, 0],
[0, 0, 1, 0, 0, 0, 0, 1, 0, 0],
[0, 0, 1, 0, 0, 0, 0, 0, 1, 0],
[0, 0, 1, 0, 0, 0, 0, 0, 0, 1],
[0, 0, 0, 1, 1, 0, 0, 0, 0, 0],
[0, 0, 0, 1, 0, 1, 0, 0, 0, 0],
[0, 0, 0, 1, 0, 0, 1, 0, 0, 0],
[0, 0, 0, 1, 0, 0, 0, 1, 0, 0],
[0, 0, 0, 1, 0, 0, 0, 0, 1, 0],
[0, 0, 0, 1, 0, 0, 0, 0, 0, 1],
[0, 0, 0, 0, 1, 1, 0, 0, 0, 0],
[0, 0, 0, 0, 1, 0, 1, 0, 0, 0],
[0, 0, 0, 0, 1, 0, 0, 1, 0, 0],
[0, 0, 0, 0, 1, 0, 0, 0, 1, 0],
[0, 0, 0, 0, 1, 0, 0, 0, 0, 1],
[0, 0, 0, 0, 0, 1, 1, 0, 0, 0],
[0, 0, 0, 0, 0, 1, 0, 1, 0, 0],
[0, 0, 0, 0, 0, 1, 0, 0, 1, 0],
[0, 0, 0, 0, 0, 1, 0, 0, 0, 1],
[0, 0, 0, 0, 0, 0, 1, 1, 0, 0],
[0, 0, 0, 0, 0, 0, 1, 0, 1, 0],
[0, 0, 0, 0, 0, 0, 1, 0, 0, 1],
[0, 0, 0, 0, 0, 0, 0, 1, 1, 0],
[0, 0, 0, 0, 0, 0, 0, 1, 0, 1],
[0, 0, 0, 0, 0, 0, 0, 0, 1, 1]],
array([21.21, 24.35, 34.57, 32.72, 36.14, 17.21, 26.41, 29.92, 32.84,
       29.09, 33.78, 40.63, 44.46, 17.03, 30.58, 32.7 , 31.19, 30.08,
       37.3 , 39.4 , 20.44, 28.34, 31.07, 36.78, 41.44, 47.97, 28.51,
       38. , 38.71, 39.24, 59.61, 29.19, 35.09, 42.18, 46.8 , 34.14,
       46.71, 48.07, 49.53, 23.46, 22.13, 24.21, 29.03, 32.51, 38.17]))

```

Let's use `least squares` to find the best solution to $M @ x == b$

```
[9]: x = la.lstsq(M,b,rcond=None)
     x[0]
```

```
[9]: array([10.63041667, 13.79291667, 13.31541667, 20.24666667, 24.32916667,
          29.46291667,  5.74916667, 14.97541667, 17.70666667, 20.11791667])
```

```
[10]: (M @ x[0],b - M @ x[0])
```

```
[10]: (array([24.42333333, 23.94583333, 30.87708333, 34.95958333, 40.09333333,
16.37958333, 25.60583333, 28.33708333, 30.74833333, 27.10833333,
34.03958333, 38.12208333, 43.25583333, 19.54208333, 28.76833333,
31.49958333, 33.91083333, 33.56208333, 37.64458333, 42.77833333,
19.06458333, 28.29083333, 31.02208333, 33.43333333, 44.57583333,
49.70958333, 25.99583333, 35.22208333, 37.95333333, 40.36458333,
53.79208333, 30.07833333, 39.30458333, 42.03583333, 44.44708333,
35.21208333, 44.43833333, 47.16958333, 49.58083333, 20.72458333,
23.45583333, 25.86708333, 32.68208333, 35.09333333, 37.82458333]),
array([-3.21333333, 0.40416667, 3.69291667, -2.23958333, -3.95333333,
0.83041667, 0.80416667, 1.58291667, 2.09166667, 1.98166667,
-0.25958333, 2.50791667, 1.20416667, -2.51208333, 1.81166667,
1.20041667, -2.72083333, -3.48208333, -0.34458333, -3.37833333,
1.37541667, 0.04916667, 0.04791667, 3.34666667, -3.13583333,
-1.73958333, 2.51416667, 2.77791667, 0.75666667, -1.12458333,
5.81791667, -0.88833333, -4.21458333, 0.14416667, 2.35291667,
-1.07208333, 2.27166667, 0.90041667, -0.05083333, 2.73541667,
-1.32583333, -1.65708333, -3.65208333, -2.58333333, 0.34541667]))
```

```
[ ]:
```